

Network Working Group
Request for Comments: 4443
Obsoletes: [2463](#)
Updates: [2780](#)
Category: Standards Track

A. Conta
Transwitch
S. Deering
Cisco Systems
M. Gupta, Ed.
Tropos Networks
March 2006

Internet Control Message Protocol (ICMPv6)
for the Internet Protocol Version 6 (IPv6) Specification

Status of This Memo

This document specifies an Internet standards track protocol for the Internet community, and requests discussion and suggestions for improvements. Please refer to the current edition of the "Internet Official Protocol Standards" (STD 1) for the standardization state and status of this protocol. Distribution of this memo is unlimited.

Copyright Notice

Copyright (C) The Internet Society (2006).

Abstract

This document describes the format of a set of control messages used in ICMPv6 (Internet Control Message Protocol). ICMPv6 is the Internet Control Message Protocol for Internet Protocol version 6 (IPv6).

Table of Contents

1. Introduction	2
2. ICMPv6 (ICMP for IPv6)	3
2.1. Message General Format	3
2.2. Message Source Address Determination	5
2.3. Message Checksum Calculation	5
2.4. Message Processing Rules	5
3. ICMPv6 Error Messages	8
3.1. Destination Unreachable Message	8
3.2. Packet Too Big Message	10
3.3. Time Exceeded Message	11
3.4. Parameter Problem Message	12
4. ICMPv6 Informational Messages	13
4.1. Echo Request Message	13
4.2. Echo Reply Message	14
5. Security Considerations	15
5.1. Authentication and Confidentiality of ICMP Messages	15
5.2. ICMP Attacks	16
6. IANA Considerations	17
6.1. Procedure for New ICMPV6 Type and Code Value Assignments ..	17
6.2. Assignments for This Document	18
7. References	19
7.1. Normative References	19
7.2. Informative References	19
8. Acknowledgements	20
Appendix A - Changes since RFC 2463.....	21

1. Introduction

The Internet Protocol version 6 (IPv6) uses the Internet Control Message Protocol (ICMP) as defined for IPv4 [RFC-792], with a number of changes. The resulting protocol is called ICMPv6 and has an IPv6 Next Header value of 58.

This document describes the format of a set of control messages used in ICMPv6. It does not describe the procedures for using these messages to achieve functions like Path MTU discovery; these procedures are described in other documents (e.g., [PMTU]). Other documents may also introduce additional ICMPv6 message types, such as Neighbor Discovery messages [IPv6-DISC], subject to the general rules for ICMPv6 messages given in Section 2 of this document.

Terminology defined in the IPv6 specification [IPv6] and the IPv6 Routing and Addressing specification [IPv6-ADDR] applies to this document as well.

This document obsoletes [RFC 2463](#) [RFC-2463] and updates [RFC 2780](#) [RFC-2780].

The key words "MUST", "MUST NOT", "REQUIRED", "SHALL", "SHALL NOT", "SHOULD", "SHOULD NOT", "RECOMMENDED", "MAY", and "OPTIONAL" in this document are to be interpreted as described in [RFC-2119].


2. ICMPv6 (ICMP for IPv6)

ICMPv6 is used by IPv6 nodes to report errors encountered in processing packets, and to perform other internet-layer functions, such as diagnostics (ICMPv6 "ping"). ICMPv6 is an integral part of IPv6, and the base protocol (all the messages and behavior required by this specification) MUST be fully implemented by every IPv6 node.

2.1. Message General Format

Every ICMPv6 message is preceded by an IPv6 header and zero or more IPv6 extension headers. The ICMPv6 header is identified by a Next Header value of 58 in the immediately preceding header. (This is different from the value used to identify ICMP for IPv4.)

The ICMPv6 messages have the following general format:



The type field indicates the type of the message. Its value determines the format of the remaining data.

The code field depends on the message type. It is used to create an additional level of message granularity.

The checksum field is used to detect data corruption in the ICMPv6 message and parts of the IPv6 header.

ICMPv6 messages are grouped into two classes: error messages and informational messages. Error messages are identified as such by a zero in the high-order bit of their message Type field values. Thus, error messages have message types from 0 to 127; informational messages have message types from 128 to 255.

This document defines the message formats for the following ICMPv6 messages:

ICMPv6 error messages:

- | | | |
|---|-------------------------|------------------------------------|
| 1 | Destination Unreachable | (see Section 3.1) |
| 2 | Packet Too Big | (see Section 3.2) |
| 3 | Time Exceeded | (see Section 3.3) |
| 4 | Parameter Problem | (see Section 3.4) |

100 Private experimentation

101 Private experimentation

127 Reserved for expansion of ICMPv6 error messages

ICMPv6 informational messages:

- | | | |
|-----|--------------|------------------------------------|
| 128 | Echo Request | (see Section 4.1) |
| 129 | Echo Reply | (see Section 4.2) |

200 Private experimentation

201 Private experimentation

255 Reserved for expansion of ICMPv6 informational messages

Type values 100, 101, 200, and 201 are reserved for private experimentation. They are not intended for general use. It is expected that multiple concurrent experiments will be done with the same type values. Any wide-scale and/or uncontrolled usage should obtain real allocations as defined in [Section 6](#).

Type values 127 and 255 are reserved for future expansion of the type value range if there is a shortage in the future. The details of this are left for future work. One possible way of doing this that would not cause any problems with current implementations is that if the type equals 127 or 255, the code field should be used for the new assignment. Existing implementations would ignore the new assignments as specified in [Section 2.4](#), (b). The new messages using these expanded type values could assign fields in the message body for its code values.

Sections [3](#) and [4](#) describe the message formats for the ICMPv6 error message types 1 through 4 and informational message types 128 and 129.

Inclusion of, at least, the start of the invoking packet is intended to allow the originator of a packet that has resulted in an ICMPv6 error message to identify the upper-layer protocol and process that sent the packet.

2.2. Message Source Address Determination

A node that originates an ICMPv6 message has to determine both the Source and Destination IPv6 Addresses in the IPv6 header before calculating the checksum. If the node has more than one unicast address, it **MUST** choose the Source Address of the message as follows:

- (a) If the message is a response to a message sent to one of the node's unicast addresses, the Source Address of the reply **MUST** be that same address.
- (b) If the message is a response to a message sent to any other address, such as
 - a multicast group address,
 - an anycast address implemented by the node, or
 - a unicast address that does not belong to the node

the Source Address of the ICMPv6 packet **MUST** be a unicast address belonging to the node. The address **SHOULD** be chosen according to the rules that would be used to select the source address for any other packet originated by the node, given the destination address of the packet. However, it **MAY** be selected in an alternative way if this would lead to a more informative choice of address reachable from the destination of the ICMPv6 packet.

2.3. Message Checksum Calculation

The checksum is the 16-bit one's complement of the one's complement sum of the entire ICMPv6 message, starting with the ICMPv6 message type field, and prepended with a "pseudo-header" of IPv6 header fields, as specified in [IPv6, [Section 8.1](#)]. The Next Header value used in the pseudo-header is 58. (The inclusion of a pseudo-header in the ICMPv6 checksum is a change from IPv4; see [[IPv6](#)] for the rationale for this change.)

For computing the checksum, the checksum field is first set to zero.

2.4. Message Processing Rules

Implementations **MUST** observe the following rules when processing ICMPv6 messages (from [[RFC-1122](#)]):

- (a) If an ICMPv6 error message of unknown type is received at its destination, it MUST be passed to the upper-layer process that originated the packet that caused the error, where this can be identified (see [Section 2.4](#), (d)).
- (b) If an ICMPv6 informational message of unknown type is received, it MUST be silently discarded.
- (c) Every ICMPv6 error message (type < 128) MUST include as much of the IPv6 offending (invoking) packet (the packet that caused the error) as possible without making the error message packet exceed the minimum IPv6 MTU [[IPv6](#)].
- (d) In cases where the internet-layer protocol is required to pass an ICMPv6 error message to the upper-layer process, the upper-layer protocol type is extracted from the original packet (contained in the body of the ICMPv6 error message) and used to select the appropriate upper-layer process to handle the error.

In cases where it is not possible to retrieve the upper-layer protocol type from the ICMPv6 message, the ICMPv6 message is silently dropped after any IPv6-layer processing. One example of such a case is an ICMPv6 message with an unusually large amount of extension headers that does not have the upper-layer protocol type due to truncation of the original packet to meet the minimum IPv6 MTU [[IPv6](#)] limit. Another example is an ICMPv6 message with an ESP extension header for which it is not possible to decrypt the original packet due to either truncation or the unavailability of the state necessary to decrypt the packet.

- (e) An ICMPv6 error message MUST NOT be originated as a result of receiving the following:
 - (e.1) An ICMPv6 error message.
 - (e.2) An ICMPv6 redirect message [[IPv6-DISC](#)].
 - (e.3) A packet destined to an IPv6 multicast address. (There are two exceptions to this rule: (1) the Packet Too Big Message ([Section 3.2](#)) to allow Path MTU discovery to work for IPv6 multicast, and (2) the Parameter Problem Message, Code 2 ([Section 3.4](#)) reporting an unrecognized IPv6 option (see [Section 4.2](#) of [[IPv6](#)]) that has the Option Type highest-order two bits set to 10).
 - (e.4) A packet sent as a link-layer multicast (the exceptions from e.3 apply to this case, too).

- (e.5) A packet sent as a link-layer broadcast (the exceptions from e.3 apply to this case, too).
- (e.6) A packet whose source address does not uniquely identify a single node -- e.g., the IPv6 Unspecified Address, an IPv6 multicast address, or an address known by the ICMP message originator to be an IPv6 anycast address.
- (f) Finally, in order to limit the bandwidth and forwarding costs incurred by originating ICMPv6 error messages, an IPv6 node **MUST** limit the rate of ICMPv6 error messages it originates. This situation may occur when a source sending a stream of erroneous packets fails to heed the resulting ICMPv6 error messages.

Rate-limiting of forwarded ICMP messages is out of scope of this specification.

A recommended method for implementing the rate-limiting function is a token bucket, limiting the average rate of transmission to N, where N can be either packets/second or a fraction of the attached link's bandwidth, but allowing up to B error messages to be transmitted in a burst, as long as the long-term average is not exceeded.

Rate-limiting mechanisms that cannot cope with bursty traffic (e.g., traceroute) are not recommended; for example, a simple timer-based implementation, allowing an error message every T milliseconds (even with low values for T), is not reasonable.

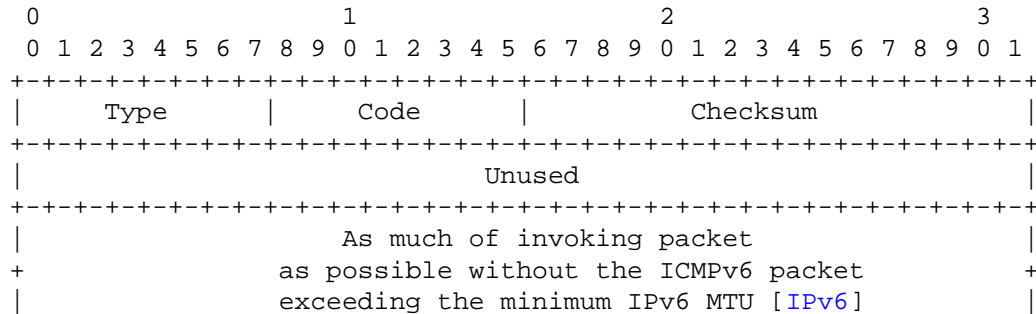
The rate-limiting parameters **SHOULD** be configurable. In the case of a token-bucket implementation, the best defaults depend on where the implementation is expected to be deployed (e.g., a high-end router vs. an embedded host). For example, in a small/mid-size device, the possible defaults could be B=10, N=10/s.

NOTE: THE RESTRICTIONS UNDER (e) AND (f) ABOVE TAKE PRECEDENCE OVER ANY REQUIREMENT ELSEWHERE IN THIS DOCUMENT FOR ORIGINATING ICMP ERROR MESSAGES.

The following sections describe the message formats for the above ICMPv6 messages.

3. ICMPv6 Error Messages

3.1. Destination Unreachable Message



IPv6 Fields:

Destination Address

Copied from the Source Address field of the invoking packet.

ICMPv6 Fields:

Type 1

Code

- 0 - No route to destination
- 1 - Communication with destination administratively prohibited
- 2 - Beyond scope of source address
- 3 - Address unreachable
- 4 - Port unreachable
- 5 - Source address failed ingress/egress policy
- 6 - Reject route to destination

Unused This field is unused for all code values. It must be initialized to zero by the originator and ignored by the receiver.

Description

A Destination Unreachable message SHOULD be generated by a router, or by the IPv6 layer in the originating node, in response to a packet that cannot be delivered to its destination address for reasons other than congestion. (An ICMPv6 message MUST NOT be generated if a packet is dropped due to congestion.)

If the reason for the failure to deliver is lack of a matching entry in the forwarding node's routing table, the Code field is set to 0.

(This error can occur only in nodes that do not hold a "default route" in their routing tables.)

If the reason for the failure to deliver is administrative prohibition (e.g., a "firewall filter"), the Code field is set to 1.

If the reason for the failure to deliver is that the destination is beyond the scope of the source address, the Code field is set to 2. This condition can occur only when the scope of the source address is smaller than the scope of the destination address (e.g., when a packet has a link-local source address and a global-scope destination address) and the packet cannot be delivered to the destination without leaving the scope of the source address.

If the reason for the failure to deliver cannot be mapped to any of other codes, the Code field is set to 3. Example of such cases are an inability to resolve the IPv6 destination address into a corresponding link address, or a link-specific problem of some sort.

One specific case in which a Destination Unreachable message is sent with a code 3 is in response to a packet received by a router from a point-to-point link, destined to an address within a subnet assigned to that same link (other than one of the receiving router's own addresses). In such a case, the packet MUST NOT be forwarded back onto the arrival link.

A destination node SHOULD originate a Destination Unreachable message with Code 4 in response to a packet for which the transport protocol (e.g., UDP) has no listener, if that transport protocol has no alternative means to inform the sender.

If the reason for the failure to deliver is that the packet with this source address is not allowed due to ingress or egress filtering policies, the Code field is set to 5.

If the reason for the failure to deliver is that the route to the destination is a reject route, the Code field is set to 6. This may occur if the router has been configured to reject all the traffic for a specific prefix.

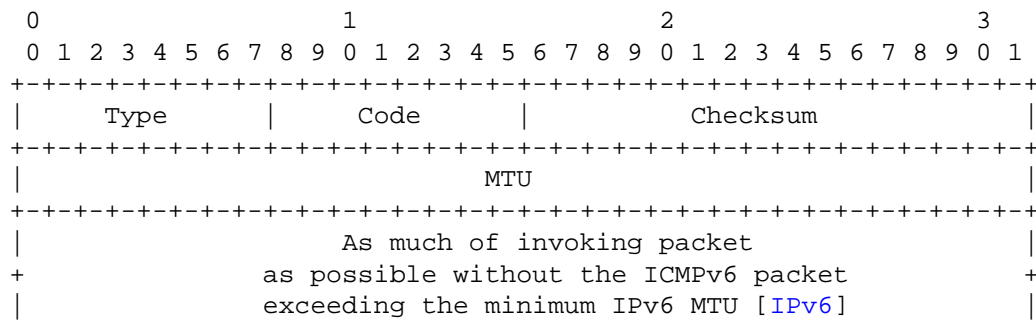
Codes 5 and 6 are more informative subsets of code 1.

For security reasons, it is recommended that implementations SHOULD allow sending of ICMP destination unreachable messages to be disabled, preferably on a per-interface basis.

Upper Layer Notification

A node receiving the ICMPv6 Destination Unreachable message MUST notify the upper-layer process if the relevant process can be identified (see [Section 2.4](#), (d)).

3.2. Packet Too Big Message



IPv6 Fields:

Destination Address

Copied from the Source Address field of the invoking packet.

ICMPv6 Fields:

Type 2

Code Set to 0 (zero) by the originator and ignored by the receiver.

MTU The Maximum Transmission Unit of the next-hop link.

Description

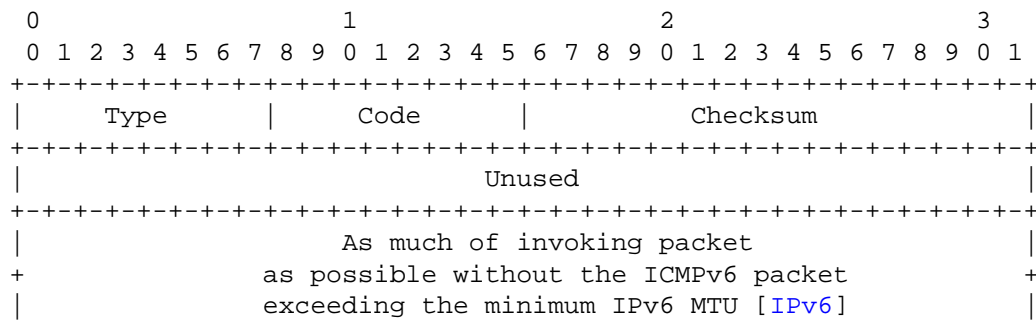
A Packet Too Big MUST be sent by a router in response to a packet that it cannot forward because the packet is larger than the MTU of the outgoing link. The information in this message is used as part of the Path MTU Discovery process [[PMTU](#)].

Originating a Packet Too Big Message makes an exception to one of the rules as to when to originate an ICMPv6 error message. Unlike other messages, it is sent in response to a packet received with an IPv6 multicast destination address, or with a link-layer multicast or link-layer broadcast address.

Upper Layer Notification

An incoming Packet Too Big message MUST be passed to the upper-layer process if the relevant process can be identified (see [Section 2.4](#), (d)).

3.3. Time Exceeded Message



IPv6 Fields:

Destination Address

Copied from the Source Address field of the invoking packet.

ICMPv6 Fields:

Type	3
Code	0 - Hop limit exceeded in transit 1 - Fragment reassembly time exceeded
Unused	This field is unused for all code values. It must be initialized to zero by the originator and ignored by the receiver.

Description

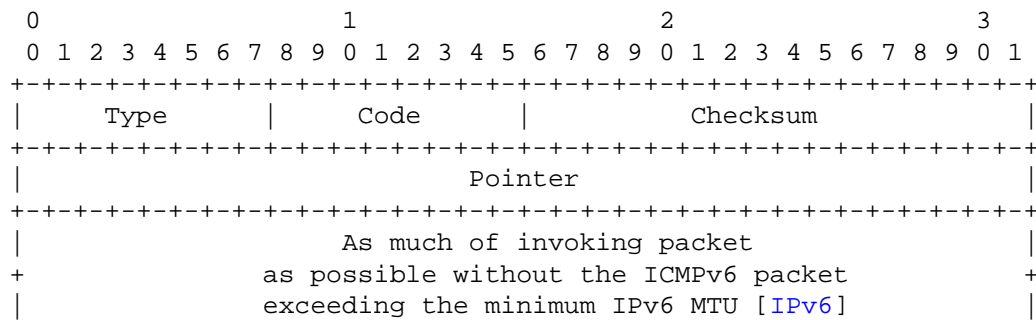
If a router receives a packet with a Hop Limit of zero, or if a router decrements a packet's Hop Limit to zero, it MUST discard the packet and originate an ICMPv6 Time Exceeded message with Code 0 to the source of the packet. This indicates either a routing loop or too small an initial Hop Limit value.

An ICMPv6 Time Exceeded message with Code 1 is used to report fragment reassembly timeout, as specified in [[IPv6](#), [Section 4.5](#)].

Upper Layer Notification

An incoming Time Exceeded message MUST be passed to the upper-layer process if the relevant process can be identified (see [Section 2.4](#), (d)).

3.4. Parameter Problem Message



IPv6 Fields:

Destination Address

Copied from the Source Address field of the invoking packet.

ICMPv6 Fields:

Type 4

Code

- 0 - Erroneous header field encountered
- 1 - Unrecognized Next Header type encountered
- 2 - Unrecognized IPv6 option encountered

Pointer Identifies the octet offset within the invoking packet where the error was detected.

The pointer will point beyond the end of the ICMPv6 packet if the field in error is beyond what can fit in the maximum size of an ICMPv6 error message.

Description

If an IPv6 node processing a packet finds a problem with a field in the IPv6 header or extension headers such that it cannot complete processing the packet, it MUST discard the packet and SHOULD originate an ICMPv6 Parameter Problem message to the packet's source, indicating the type and location of the problem.

Codes 1 and 2 are more informative subsets of Code 0.


The pointer identifies the octet of the original packet's header where the error was detected. For example, an ICMPv6 message with a Type field of 4, Code field of 1, and Pointer field of 40 would indicate that the IPv6 extension header following the IPv6 header of the original packet holds an unrecognized Next Header field value.

Upper Layer Notification

A node receiving this ICMPv6 message **MUST** notify the upper-layer process if the relevant process can be identified (see [Section 2.4](#), (d)).

4. ICMPv6 Informational Messages

4.1. Echo Request Message



IPv6 Fields:

Destination Address

Any legal IPv6 address.

ICMPv6 Fields:

Type 128

Code 0

Identifier An identifier to aid in matching Echo Replies to this Echo Request. May be zero.

Sequence Number

A sequence number to aid in matching Echo Replies to this Echo Request. May be zero.


Data Zero or more octets of arbitrary data.

Description

Every node **MUST** implement an ICMPv6 Echo responder function that receives Echo Requests and originates corresponding Echo Replies. A node **SHOULD** also implement an application-layer interface for originating Echo Requests and receiving Echo Replies, for diagnostic purposes.

Upper Layer Notification

Echo Request messages **MAY** be passed to processes receiving ICMP messages.

4.2. Echo Reply Message**IPv6 Fields:****Destination Address**

Copied from the Source Address field of the invoking Echo Request packet.

ICMPv6 Fields:

Type 129

Code 0

Identifier The identifier from the invoking Echo Request message.

Sequence Number

The sequence number from the invoking Echo Request message.

Data The data from the invoking Echo Request message.

Description

Every node MUST implement an ICMPv6 Echo responder function that receives Echo Requests and originates corresponding Echo Replies. A node SHOULD also implement an application-layer interface for originating Echo Requests and receiving Echo Replies, for diagnostic purposes.

The source address of an Echo Reply sent in response to a unicast Echo Request message MUST be the same as the destination address of that Echo Request message.

An Echo Reply SHOULD be sent in response to an Echo Request message sent to an IPv6 multicast or anycast address. In this case, the source address of the reply MUST be a unicast address belonging to the interface on which the Echo Request message was received.

The data received in the ICMPv6 Echo Request message MUST be returned entirely and unmodified in the ICMPv6 Echo Reply message.

Upper Layer Notification

Echo Reply messages MUST be passed to the process that originated an Echo Request message. An Echo Reply message MAY be passed to processes that did not originate the Echo Request message.

Note that there is no limitation on the amount of data that can be put in Echo Request and Echo Reply Messages.

5. Security Considerations

5.1. Authentication and Confidentiality of ICMP Messages

ICMP protocol packet exchanges can be authenticated using the IP Authentication Header [IPv6-AUTH] or IP Encapsulating Security Payload Header [IPv6-ESP]. Confidentiality for the ICMP protocol packet exchanges can be achieved using the IP Encapsulating Security Payload Header [IPv6-ESP].

[SEC-ARCH] describes the IPsec handling of ICMP traffic in detail.

5.2. ICMP Attacks

ICMP messages may be subject to various attacks. A complete discussion can be found in the IP Security Architecture [IPv6-SA]. A brief discussion of these attacks and their prevention follows:

1. ICMP messages may be subject to actions intended to cause the receiver to believe the message came from a different source from that of the message originator. The protection against this attack can be achieved by applying the IPv6 Authentication mechanism [IPv6-AUTH] to the ICMP message.
2. ICMP messages may be subject to actions intended to cause the message or the reply to it to go to a destination different from that of the message originator's intention. The protection against this attack can be achieved by using the Authentication Header [IPv6-AUTH] or the Encapsulating Security Payload Header [IPv6-ESP]. The Authentication Header provides the protection against change for the source and the destination address of the IP packet. The Encapsulating Security Payload Header does not provide this protection, but the ICMP checksum calculation includes the source and the destination addresses, and the Encapsulating Security Payload Header protects the checksum. Therefore, the combination of ICMP checksum and the Encapsulating Security Payload Header provides protection against this attack. The protection provided by the Encapsulating Security Payload Header will not be as strong as the protection provided by the Authentication Header.
3. ICMP messages may be subject to changes in the message fields, or payload. The authentication [IPv6-AUTH] or encryption [IPv6-ESP] of the ICMP message protects against such actions.
4. ICMP messages may be used to attempt denial-of-service attacks by sending back to back erroneous IP packets. An implementation that correctly followed Section 2.4, paragraph (f), of this specification, would be protected by the ICMP error rate limiting mechanism.
5. The exception number 2 of rule e.3 in Section 2.4 gives a malicious node the opportunity to cause a denial-of-service attack to a multicast source. A malicious node can send a multicast packet with an unknown destination option marked as mandatory, with the IPv6 source address of a valid multicast source. A large number of destination nodes will send an ICMP Parameter Problem Message to the multicast source, causing a denial-of-service attack. The way multicast traffic is forwarded by the multicast routers requires that the malicious node be part of the correct

multicast path, i.e., near to the multicast source. This attack can only be avoided by securing the multicast traffic. The multicast source should be careful while sending multicast traffic with the destination options marked as mandatory, because they can cause a denial-of-service attack to themselves if the destination option is unknown to a large number of destinations.

6. As the ICMP messages are passed to the upper-layer processes, it is possible to perform attacks on the upper layer protocols (e.g., TCP) with ICMP [[TCP-attack](#)]. It is recommended that the upper layers perform some form of validation of ICMP messages (using the information contained in the payload of the ICMP message) before acting upon them. The actual validation checks are specific to the upper layers and are out of the scope of this specification. Protecting the upper layer with IPsec mitigates these attacks.

ICMP error messages signal network error conditions that were encountered while processing an internet datagram. Depending on the particular scenario, the error conditions being reported might or might not get solved in the near term. Therefore, reaction to ICMP error messages may depend not only on the error type and code but also on other factors, such as the time at which the error messages are received, previous knowledge of the network error conditions being reported, and knowledge of the network scenario in which the receiving host is operating.

6. IANA Considerations

6.1. Procedure for New ICMPV6 Type and Code Value Assignments

The IPv6 ICMP header defined in this document contains the following fields that carry values assigned from IANA-managed name spaces: Type and Code. Code field values are defined relative to a specific Type value.

Values for the IPv6 ICMP Type fields are allocated using the following procedure:

1. The IANA should allocate and permanently register new ICMPv6 type codes from IETF RFC publication. This is for all RFC types, including standards track, informational, and experimental status, that originate from the IETF and have been approved by the IESG for publication.
2. IETF working groups with working group consensus and area director approval can request reclaimable ICMPV6 type code assignments from the IANA. The IANA will tag the values as "reclaimable in future".

The "reclaimable in the future" tag will be removed when an RFC is published that documents the protocol as defined in 1. This will make the assignment permanent and update the reference on the IANA web pages.

At the point where the ICMPv6 type values are 85% assigned, the IETF will review the assignments tagged "reclaimable in the future" and inform the IANA which ones should be reclaimed and reassigned.

3. Requests for new ICMPv6 type value assignments from outside the IETF are only made through the publication of an IETF document, per 1 above. Note also that documents published as "RFC Editor contributions" [RFC-3978] are not considered IETF documents.

The assignment of new Code values for the Type values defined in this document require standards action or IESG approval. The policy for assigning Code values for new IPv6 ICMP Types not defined in this document should be defined in the document defining the new Type values.

6.2. Assignments for This Document

The following has updated assignments located at:

<http://www.iana.org/assignments/icmpv6-parameters>

The IANA has reassigned ICMPv6 type 1 "Destination Unreachable" code 2, which was unassigned in [RFC-2463], to:

2 - Beyond scope of source address

The IANA has assigned the following two new codes values for ICMPv6 type 1 "Destination Unreachable":

5 - Source address failed ingress/egress policy
6 - Reject route to destination

The IANA has assigned the following new type values:

100 Private experimentation
101 Private experimentation

127 Reserved for expansion of ICMPv6 error messages

200 Private experimentation
201 Private experimentation

255 Reserved for expansion of ICMPv6 informational messages

7. References

7.1. Normative References

- [IPv6] Deering, S. and R. Hinden, "Internet Protocol, Version 6 (IPv6) Specification", [RFC 2460](#), December 1998.
- [IPv6-DISC] Narten, T., Nordmark, E., and W. Simpson, "Neighbor Discovery for IP Version 6 (IPv6)", [RFC 2461](#), December 1998.
- [RFC-792] Postel, J., "Internet Control Message Protocol", STD 5, [RFC 792](#), September 1981.
- [RFC-2463] Conta, A. and S. Deering, "Internet Control Message Protocol (ICMPv6) for the Internet Protocol Version 6 (IPv6) Specification", [RFC 2463](#), December 1998.
- [RFC-1122] Braden, R., "Requirements for Internet Hosts - Communication Layers", STD 3, [RFC 1122](#), October 1989.
- [RFC-2119] Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", [BCP 14](#), [RFC 2119](#), March 1997.
- [RFC-3978] Bradner, S., "IETF Rights in Contributions", [BCP 78](#), [RFC 3978](#), March 2005.

7.2. Informative References

- [RFC-2780] Bradner, S. and V. Paxson, "IANA Allocation Guidelines For Values In the Internet Protocol and Related Headers", [BCP 37](#), [RFC 2780](#), March 2000.
- [IPv6-ADDR] Hinden, R. and S. Deering, "Internet Protocol Version 6 (IPv6) Addressing Architecture", [RFC 3513](#), April 2003.
- [PMTU] McCann, J., Deering, S., and J. Mogul, "Path MTU Discovery for IP version 6", [RFC 1981](#), August 1996.
- [IPv6-SA] Kent, S. and R. Atkinson, "Security Architecture for the Internet Protocol", [RFC 2401](#), November 1998.
- [IPv6-AUTH] Kent, S., "IP Authentication Header", [RFC 4302](#), December 2005.

- [IPv6-ESP] Kent, S., "IP Encapsulating Security Payload (ESP)", [RFC 4203](#), December 2005.
- [SEC-ARCH] Kent, S. and K. Seo, "Security Architecture for the Internet Protocol", [RFC 4301](#), December 2005.
- [TCP-attack] Gont, F., "[ICMP attacks against TCP](#)", Work in Progress.

8. Acknowledgements

The document is derived from previous ICMP documents of the SIPP and IPng working group.

The IPng working group, and particularly Robert Elz, Jim Bound, Bill Simpson, Thomas Narten, Charlie Lynn, Bill Fink, Scott Bradner, Dimitri Haskin, Bob Hinden, Jun-ichiro Itojun Hagino, Tatuya Jinmei, Brian Zill, Pekka Savola, Fred Templin, and Elwyn Davies (in chronological order) provided extensive review information and feedback.

Bob Hinden was the document editor for this document.

Appendix A - Changes since RFC 2463

The following changes were made from RFC 2463:

- Edited the Abstract to make it a little more elaborate.
- Corrected typos in Section 2.4, where references to sub-bullet e.2 were supposed to be references to e.3.
- Removed the Timer-based and the Bandwidth-based methods from the example rate-limiting mechanism for ICMP error messages. Added Token-bucket based method.
- Added specification that all ICMP error messages shall have exactly 32 bits of type-specific data, so that receivers can reliably find the embedded invoking packet even when they don't recognize the ICMP message Type.
- In the description of Destination Unreachable messages, Code 3, added rule prohibiting forwarding of packets back onto point-to-point links from which they were received, if their destination addresses belong to the link itself ("anti-ping-ponging" rule).
- Added description of Time Exceeded Code 1 (fragment reassembly timeout).
- Added "beyond scope of source address", "source address failed ingress/egress policy", and "reject route to destination" messages to the family of "unreachable destination" type ICMP error messages (Section 3.1).
- Reserved some ICMP type values for experimentation.
- Added a NOTE in Section 2.4 that specifies ICMP message processing rules precedence.
- Added ICMP REDIRECT to the list in Section 2.4, (e) of cases in which ICMP error messages are not to be generated.
- Made minor editorial changes in Section 2.3 on checksum calculation, and in Section 5.2.
- Clarified in Section 4.2, regarding the Echo Reply Message; the source address of an Echo Reply to an anycast Echo Request should be a unicast address, as in the case of multicast.

- Revised the Security Considerations section. Added the use of the Encapsulating Security Payload Header for authentication. Changed the requirement of an option of "not allowing unauthenticated ICMP messages" to MAY from SHOULD.
- Added a new attack in the list of possible ICMP attacks in [Section 5.2](#).
- Separated References into Normative and Informative.
- Added reference to [RFC 2780](#) "IANA Allocation Guidelines For Values In the Internet Protocol and Related Headers". Also added a note that this document updates [RFC 2780](#).
- Added a procedure for new ICMPv6 Type and Code value assignments in the IANA Considerations section.
- Replaced word "send" with "originate" to make it clear that ICMP packets being forwarded are out of scope of this specification.
- Changed the ESP and AH references to the updated ESP and AH documents.
- Added reference to the updated IPsec Security Architecture document.
- Added a SHOULD requirement for allowing the sending of ICMP destination unreachable messages to be disabled.
- Simplified the source address selection of the ICMPv6 packet.
- Reorganized the General Message Format ([Section 2.1](#)).
- Removed the general packet format from [Section 2.1](#). It refers to [Sections 3](#) and [4](#) for packet formats now.
- Added text about attacks to the transport protocols that could potentially be caused by ICMP.

Authors' Addresses

Alex Conta
Transwitch Corporation
3 Enterprise Drive
Shelton, CT 06484
USA

EMail: aconta@txc.com

Stephen Deering
Cisco Systems, Inc.
170 West Tasman Drive
San Jose, CA 95134-1706
USA

Mukesh Gupta, Ed.
Tropos Networks
555 Del Rey Avenue
Sunnyvale, CA 94085

Phone: +1 408-331-6889
EMail: mukesh.gupta@tropos.com

Full Copyright Statement

Copyright (C) The Internet Society (2006).

This document is subject to the rights, licenses and restrictions contained in [BCP 78](#), and except as set forth therein, the authors retain all their rights.

This document and the information contained herein are provided on an "AS IS" basis and THE CONTRIBUTOR, THE ORGANIZATION HE/SHE REPRESENTS OR IS SPONSORED BY (IF ANY), THE INTERNET SOCIETY AND THE INTERNET ENGINEERING TASK FORCE DISCLAIM ALL WARRANTIES, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OR ANY IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Intellectual Property

The IETF takes no position regarding the validity or scope of any Intellectual Property Rights or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; nor does it represent that it has made any independent effort to identify any such rights. Information on the procedures with respect to rights in RFC documents can be found in [BCP 78](#) and [BCP 79](#).

Copies of IPR disclosures made to the IETF Secretariat and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the IETF on-line IPR repository at <http://www.ietf.org/ipr>.

The IETF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights that may cover technology that may be required to implement this standard. Please address the information to the IETF at ietf-ipr@ietf.org.

Acknowledgement

Funding for the RFC Editor function is provided by the IETF Administrative Support Activity (IASA).